

TEKLRN

HADOOP LEVEL 16

1. Managing HDFS
 - Where to write Data
 - Using Balancer
 - When to rebalance
2. MapReduce Management
 - Command line job Management
 - Job Priorities and scheduling
3. Changing job Priorities and killing Jobs
 - Alternative scheduler
 - Capacity Scheduler
 - Fair Scheduler
 - Enabling Alternative schedulers
 - When to use alternative schedulers
4. Scaling
 - Adding Capacity to a local Hadoop Cluster
 - Adding Capacity to an EMR job flow
 - Expanding a running Job flow