- **Training & evaluation with the built-in methods**
- **API Overview a first end-to-end example**
- **use the MNIST dataset as NumPy arrays, in order to demonstrate how to use optimizers, losses, and metrics.**
- **What and end-to-end workflow looks like**
- **Demonstration with the training configuration (optimizer, loss, metrics):**
- **Usage of fit() method for Training the model**
- **Usage of evaluate()**
- **The compile() method: specifying a loss, metrics, and an optimizer**
- **Custom losses**
- **Custom metrics**
- **Handling losses and metrics that don't fit the standard signature**
- **Automatically setting apart a validation holdout set**
- **Training & evaluation from tf.data Datasets**
- **Using a validation dataset**
- **Using a keras.utils.Sequence object as input**
- **Passing data to multi-input and multioutput models**
- **Writing your own callback.**
- **Checkpointing models.**
- **Using learning rate schedules**
- **Visualizing loss and metrics during training**