

- **Customizing what happens in fit()**
- **Introduction**
- **Setup**
- **A first simple example**
- **Going lower-level**
- **a lower-level example, that only uses compile() to configure the optimizer**
- **call reset_states() on your metrics after each epoch, or between training and evaluation.**
- **Supporting sample_weight & class_weight**
- **Providing your own evaluation step**
- **What if you want to do the same for calls to model.evaluate()?**
- **Wrapping up: an end-to-end GAN example with the conditions below:**
- **A generator network meant to generate 28x28x1 images.**
- **A discriminator network meant to classify 28x28x1 images into two classes ("fake" and "real").**
- **One optimizer for each.**
- **A loss function to train the discriminator**