

TEKLRN

KERAS LEVEL 27

1. Multi-GPU and distributed training
 - Introduction
 - Setup
 - Data parallelism
 - Model parallelism
2. Single-host, multi-device synchronous training
 - How it works
 - How to use it
3. Using callbacks to ensure fault tolerance
4. tf.data performance tips
5. Note about dataset batching
 - Calling `dataset.cache()`
 - Calling `dataset.prefetch(buffer_size)`
6. Multi-worker distributed synchronous training
 - How it works
 - How to use it
 - Setting up a cluster
 - Setting up the `TF_CONFIG` environment variable
 - Running code on each worker
 - Example: code running in a multi-worker setup