

- **Multi-GPU and distributed training**
- **Introduction**
- **Setup**
- **Data parallelism**
- **Model parallelism**
- **Single-host, multi-device synchronous training**
- **Using callbacks to ensure fault tolerance**
- **tf.data performance tips**
- **Note about dataset batching**
- **Calling dataset.cache()**
- **Calling dataset.prefetch(buffer_size)**
- **Multi-worker distributed synchronous training**
- **Setting up a cluster**
- **Setting up the TF_CONFIG environment variable**
- **Running code on each worker**
- **Example: code running in a multi-worker setup**