

PyTorch LEVEL 5

1. CUDA SEMANTICS
2. TensorFloat-32(TF32) on Ampere devices
3. Asynchronous execution
4. CUDA streams
5. Memory management
6. cuFFT plan cache
7. Best practices
8. Device-agnostic code
9. Use pinned memory buffers
10. Use `nn.parallel.DistributedDataParallel` instead of `multiprocessing` or `nn.DataParallel`